

Analysis of Vettius Valens

Definitions

```
library(ggplot2)
library(dplyr)

##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(cowplot)

##
## *****
## Note: As of version 1.0.0, cowplot does not change the
##   default ggplot2 theme anymore. To recover the previous
##   behavior, execute:
##   theme_set(theme_cowplot())
## *****

library(tidyr)
library(broman)

# Order of bodies
ORDER_OF_BODIES = c('Saturn',
                    'Jupiter',
                    'Mars',
                    'Sun',
                    'Venus',
                    'Mercury',
                    'Moon')
names(ORDER_OF_BODIES) <- c("Sa", "J", "Mar", "Su", "V", "Mer", "Moo")
# For plots
theme_basic <- function () {
  theme_bw(base_size=12) %>%replace%
  theme(
    axis.text=element_text(colour="black")
  ) %>%replace%
  theme(
    panel.grid=element_blank()
```

```
)
}
```

Singles

We read in the sentiments of the singles.

```
singles.sentiments <- read.csv('../data/singles-qualities.csv',
                                header=T,
                                stringsAsFactors = F)

# Get in useful format
singles.sentiments.df <- singles.sentiments %>% group_by(PLANET, SENTIMENT) %>%
  summarise(count=n()) %>%
  mutate(total=sum(count),
          prop=count/total)
# remove 'total' variable
singles.sentiments.df$total <- NULL
# Make sure list is complete (some sentiments are missing where none listed in input data)
full_data <- expand.grid(PLANET=singles.sentiments.df$PLANET, SENTIMENT=singles.sentiments.df$SENTIMENT)
singles.sentiments.df <- unique(left_join(tbl_df(full_data), singles.sentiments.df))

## Joining, by = c("PLANET", "SENTIMENT")

## Warning: Column `PLANET` joining factor and character vector, coercing into
## character vector

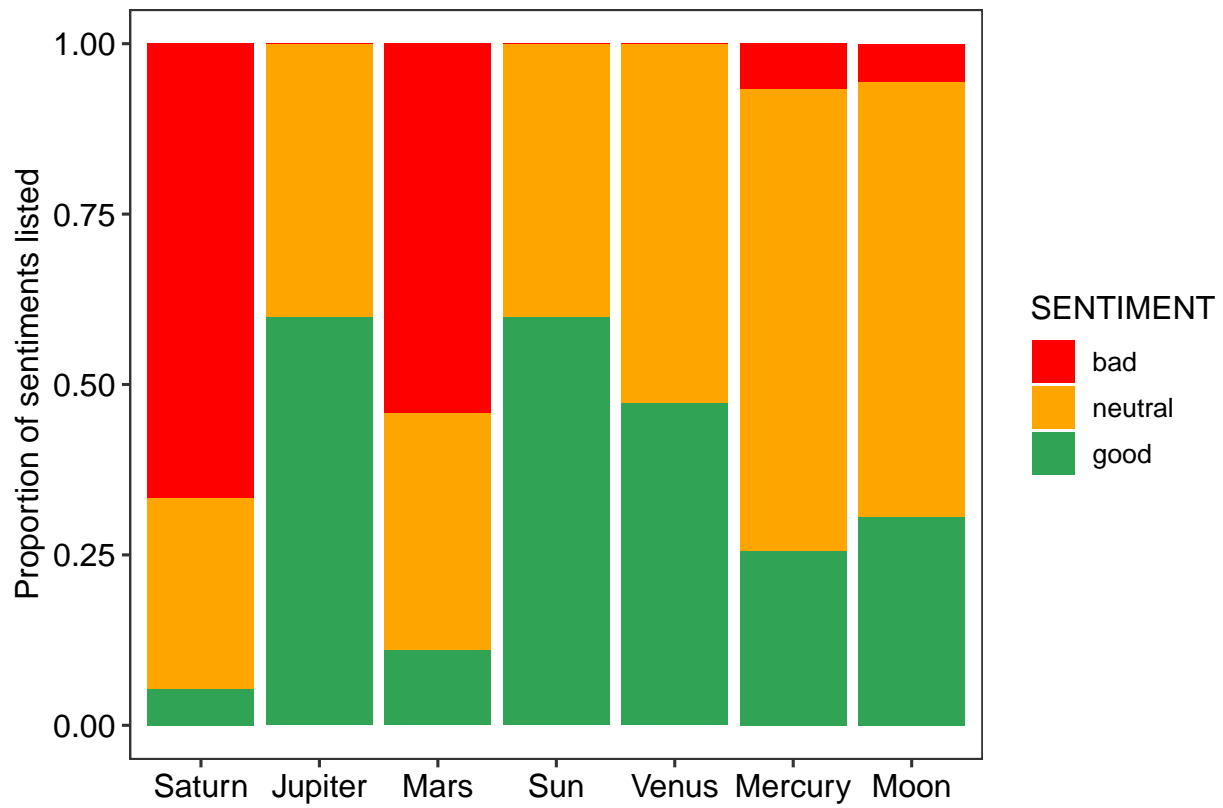
## Warning: Column `SENTIMENT` joining factor and character vector, coercing into
## character vector

singles.sentiments.df$count[is.na(singles.sentiments.df$count)] <- 0
singles.sentiments.df$prop[is.na(singles.sentiments.df$prop)] <- 0
# Order planets and sentiments
singles.sentiments.df$PLANET <- ordered(singles.sentiments.df$PLANET,
                                       levels=ORDER_OF_BODIES)
singles.sentiments.df$SENTIMENT <- ordered(singles.sentiments.df$SENTIMENT,
                                       levels=c("bad", "neutral", "good"))

# Overall sentiment index
singles.overall.sentiments <- singles.sentiments.df %>% group_by(PLANET) %>%
  summarise(sentiment=prop[which(SENTIMENT=="good")] - prop[which(SENTIMENT=="bad")])
```


We then plot the sentiments of the singles.

```
# Plot them
ggplot(singles.sentiments.df, aes(PLANET, prop, fill=SENTIMENT))+
  geom_bar(stat="identity")+
  theme_basic()+
  xlab("")+
  scale_fill_manual(values=c("red", "orange", "#31a354"))+
  ylab("Proportion of sentiments listed")
```



Also plot overall sentiment index (good - bad).

```
ggplot(singles.overall.sentiments, aes(PLANET, sentiment))+  
  geom_bar(stat="identity", fill="black")+  
  theme_basic()+  
  xlab("")+  
  ylab("Sentiment index")
```



Doubles

We read in the sentiments of the double conjunctions (e.g. Saturn and Jupiter).

```
doubles.sentiments <- read.csv('../data/doubles-qualities.csv',
                                header=T,
                                stringsAsFactors = F)

# Get in useful format
doubles.sentiments.df <- doubles.sentiments %>% group_by(DOUBLE, SENTIMENT) %>%
  summarise(count=n()) %>%
  mutate(total=sum(count),
         prop=count/total)
# remove 'total' variable
doubles.sentiments.df$total <- NULL
# Make sure list is complete (some sentiments are missing where none listed in input data)
full_data <- expand.grid(DOUBLE=doubles.sentiments.df$DOUBLE, SENTIMENT=doubles.sentiments.df$SENTIMENT)
doubles.sentiments.df <- unique(left_join(tbl_df(full_data),doubles.sentiments.df))

## Joining, by = c("DOUBLE", "SENTIMENT")

## Warning: Column `DOUBLE` joining factor and character vector, coercing into
## character vector

## Warning: Column `SENTIMENT` joining factor and character vector, coercing into
## character vector

doubles.sentiments.df$count[is.na(doubles.sentiments.df$count)] <- 0
doubles.sentiments.df$prop[is.na(doubles.sentiments.df$prop)] <- 0

# Order planets and sentiments
```

```

doubles.sentiments.df$body.1 <- sapply(stringr::str_split(doubles.sentiments.df$DOUBLE, pattern=" "),
  function(x) x[1])
doubles.sentiments.df$body.1 <- ordered(ORDER_OF_BODIES[doubles.sentiments.df$body.1],
  levels=ORDER_OF_BODIES)
doubles.sentiments.df$body.2 <- sapply(stringr::str_split(doubles.sentiments.df$DOUBLE, pattern=" "),
  function(x) x[2])
doubles.sentiments.df$body.2 <- ordered(ORDER_OF_BODIES[doubles.sentiments.df$body.2],
  levels=ORDER_OF_BODIES)
doubles.sentiments.df$bodies.sorted <- sapply(1:nrow(doubles.sentiments.df),
  function(x)
    paste(as.character(sort(unlist(c(doubles.sentiments.df[x,"body.1"],
                                     doubles.sentiments.df[x,"body.2"]
                                     collapse=" ")))
    # Make body 1 and 2 be in order of bodies as expected
doubles.sentiments.df$body.1 <- gsub(".*", "", doubles.sentiments.df$bodies.sorted)
doubles.sentiments.df$body.2 <- gsub(".* ", "", doubles.sentiments.df$bodies.sorted)
doubles.sentiments.df$body.1 <- ordered(doubles.sentiments.df$body.1,
  levels=ORDER_OF_BODIES)
doubles.sentiments.df$body.2 <- ordered(doubles.sentiments.df$body.2,
  levels=ORDER_OF_BODIES)
doubles.sentiments.df$order.body.string <- paste0(as.numeric(doubles.sentiments.df$body.1),
  as.numeric(doubles.sentiments.df$body.2))
# Order bodies again (this is hacky but it works)
doubles.sentiments.df$bodies.sorted <- ordered(doubles.sentiments.df$bodies.sorted,
  levels=unique(doubles.sentiments.df$bodies.sorted[order(d
doubles.sentiments.df$SENTIMENT <- ordered(doubles.sentiments.df$SENTIMENT,
  levels=c("bad", "good"))
# Also make a dataframe with an overall sentiment index
doubles.overall.sentiments <- doubles.sentiments.df %>% group_by(body.1, body.2, bodies.sorted, order.b
  summarise(sentiment=prop[which(SENTIMENT=="good")]-prop[which(SENTIMENT=="bad")])


```

We then plot the sentiments of the doubles.

```

# Plot them
ggplot(doubles.sentiments.df, aes(bodies.sorted, prop, fill=SENTIMENT))+
  geom_bar(stat="identity")+
  theme_basic()+
  xlab("")+
  scale_fill_manual(values=c("red", "#31a354"))+
  ylab("Proportion of sentiments listed")+
  theme(axis.text.x=element_text(angle=45, hjust=1))

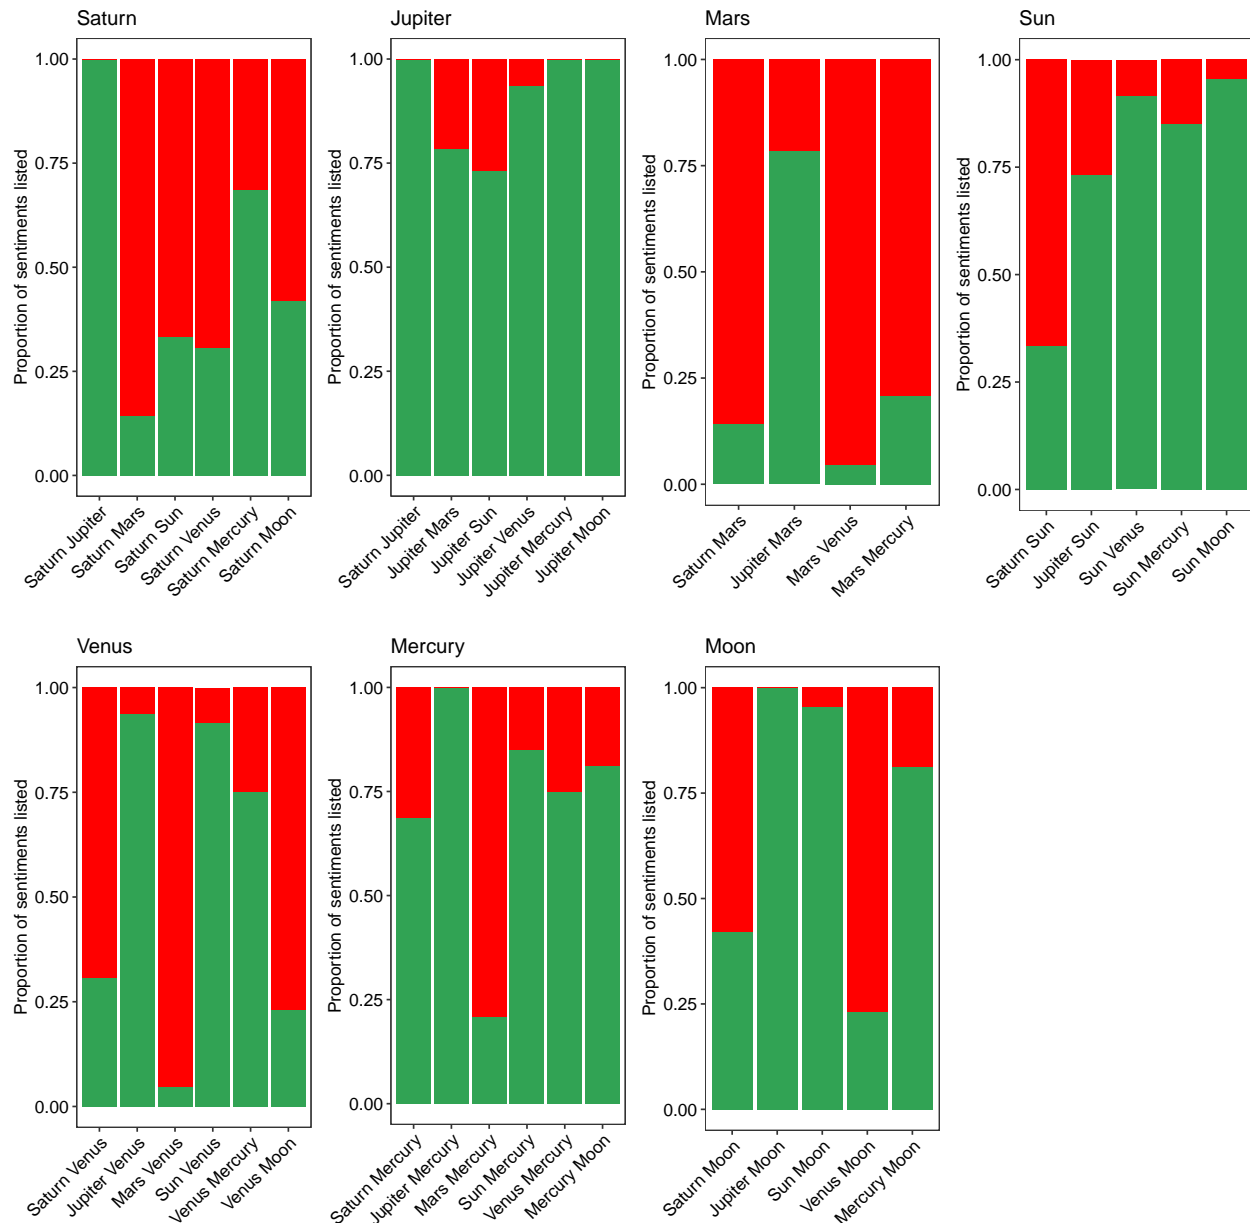
```



A function to plot all the doubles involving a particular planet.

```
plotDoublesWithPlanet <- function(planet){
  # Subset to only doubles involving planet
  local.planet.df <- doubles.sentiments.df[grep(planet, doubles.sentiments.df$bodies.sorted),]
  ggplot(local.planet.df, aes(bodies.sorted, prop, fill=SENTIMENT))+
    geom_bar(stat="identity")+
    theme_basic()+
    xlab("")+
    scale_fill_manual(values=c("red", "#31a354"))+
    ylab("Proportion of sentiments listed")+
    theme(axis.text.x=element_text(angle=45, hjust=1))+
    ggtitle(planet)+
    theme(legend.position = "none")
}

# Make all these plots and combine them
p.doubles.saturn <- plotDoublesWithPlanet("Saturn")
p.doubles.jupiter <- plotDoublesWithPlanet("Jupiter")
p.doubles.mars <- plotDoublesWithPlanet("Mars")
p.doubles.sun <- plotDoublesWithPlanet("Sun")
p.doubles.venus <- plotDoublesWithPlanet("Venus")
p.doubles.mercury <- plotDoublesWithPlanet("Mercury")
p.doubles.moon <- plotDoublesWithPlanet("Moon")
cowplot::plot_grid(p.doubles.saturn, p.doubles.jupiter,
  p.doubles.mars, p.doubles.sun,
  p.doubles.venus, p.doubles.mercury,
  p.doubles.moon, nrow=2)
```



Note that Valens misses out two doubles which involve Mars: Mars + Sun, and Mars + Moon.

Predicting doubles from singles

If we know the sentiment proportions associated with single planets, can we predict the sentiment of the doubles?

```
# getSingleScores <- function(double, sentiment="good", mean=TRUE, singles=singles.sentiments.df){
#   double.bodies <- ordered(unlist(stringr::str_split(double, pattern=" "),
#                               levels=ORDER_OF_BODIES)
#   prop <- 0
#   for (body in double.bodies){
#     new.prop <- as.numeric(singles[which(singles$PLANET==body &
#                                           singles$SENTIMENT==sentiment), "prop"] )
#     prop <- prop + new.prop
#   }
# }
```

```

#   }
#   if (mean==FALSE){
#     return(prop)
#   }
#   else{
#     return(prop/2)
#   }
# }
getSingleScores <- function(conjunction, mean=TRUE, singles=singles.overall.sentiments){
  conjunction.bodies <- ordered(unlist(stringr::str_split(conjunction, pattern=" "),
    levels=ORDER_OF_BODIES)

  prop <- 0
  for (body in conjunction.bodies){
    new.prop <- as.numeric(singles[which(singles$PLANET==body), "sentiment"])
    prop <- prop + new.prop
  }
  if (mean==FALSE){
    return(prop)
  }
  else{
    return(prop/length(conjunction.bodies))
  }
}

doubles.overall.sentiments$mean.single.sentiments <- sapply(doubles.overall.sentiments$bodies.sorted,
  function(x) getSingleScores(x, mean=TRUE))


# Add correlation score
spearman.cor <- cor.test(doubles.overall.sentiments$mean.single.sentiments, doubles.overall.sentiments$

## Warning in cor.test.default(doubles.overall.sentiments$mean.single.sentiments, :
## Cannot compute exact p-value with ties

p.sentiments.doubles.singles <- ggplot(doubles.overall.sentiments, aes(mean.single.sentiments, sentiment))
  stat_smooth(method="lm", se=FALSE, colour="red")+
  geom_point(colour="black")+
  theme_basic()+
  ggrepel::geom_text_repel(aes(label=bodies.sorted))+
  coord_fixed()+
  xlim(c(-1, 1))+
  ylim(c(-1,1))+
  geom_hline(yintercept = 0, linetype='dashed')+
  geom_vline(xintercept = 0, linetype='dashed')+
  xlab("Mean sentiment index of single planets")+
  ylab("Sentiment index of double")+
  annotate(geom="label", label=paste0("Spearman's rho = ", myround(spearman.cor$estimate, 3), "\n",
    "p = ", myround(spearman.cor$p.value, 3)), x=1, y=-0.75, hjust=1)

p.sentiments.doubles.singles

```

We see that there is a strong correlation between the mean sentiment index for the component singles and the sentiment index of the double conjunction.

Double conjunction occurrences

Read in the frequencies of the double conjunctions.

```
double.occurrences <- read.csv('../data/0-CE-200-CE-double-occurrence-per-year.csv', header=F, stringsAsFactors=F)
colnames(double.occurrences) <- c("year", "double", "occurrence")
double.occurrences.median <- double.occurrences %>% group_by(double) %>%
  summarise(median=median(occurrence))
order.of.doubles <- double.occurrences.median$double[order(double.occurrences.median$median, decreasing=T)]
double.occurrences.median$year <- 160
double.occurrences.median$y <- 11
double.occurrences.median$double <- ordered(double.occurrences.median$double,
```


```

                                levels=order.of.doubles)
double.occurrences.median$median.plot <- myround(double.occurrences.median$median, 1)
double.occurrences$double <- ordered(double.occurrences$double,
                                levels=order.of.doubles)

# Make plot
occurrences.plot <- ggplot(double.occurrences, aes(year, occurrence, group=double))+
  geom_line()+
  theme_bw()+
  xlab("Year (CE)")+
  ylab("Z-codes in year containing double (%)")+
  facet_wrap(~double, ncol=7)+
  theme(panel.grid = element_blank())+
  ylim(c(0,100))+
  xlim(c(120, 175))+
  theme(strip.text = element_text(size=8))+
  geom_text(data=double.occurrences.median,
            aes(year, 90, label=paste0(median.plot, "%"), group=double),
            hjust=0, size=3)
occurrences.plot


```

```
## Warning: Removed 3003 rows containing missing values (geom_path).
```



Does median occurrence relate to mean sentiment index?

```
double.occurrences.median$bodies.sorted <- ordered(gsub("-", " ", double.occurrences.median$double),
                                                    levels=levels(doubles.overall.sentiments$bodies.sorted))
median.occurrences.doubles <- double.occurrences.median$median
names(median.occurrences.doubles) <- double.occurrences.median$bodies.sorted
# Add to sentiment
doubles.overall.sentiments$median.occurrence <- median.occurrences.doubles[doubles.overall.sentiments$bodies.sorted]
ggplot(doubles.overall.sentiments, aes(median.occurrence, sentiment))+
  geom_point()+
  theme_basic()
```



Not obviously. But perhaps there is a linear model term we can use.

```
summary(lm(sentiment ~ mean.single.sentiments + median.occurrence, data=doubles.overall.sentiments))
```

```
##
## Call:
## lm(formula = sentiment ~ mean.single.sentiments + median.occurrence,
##     data = doubles.overall.sentiments)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.07939 -0.35603  0.05497  0.33136  0.95689
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.0565613   0.1916160    0.295  0.77165
## mean.single.sentiments 1.3581772   0.4166253    3.260  0.00492 **
## median.occurrence   -0.0008909   0.0114380   -0.078  0.93888
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5439 on 16 degrees of freedom
## Multiple R-squared:  0.4017, Adjusted R-squared:  0.327
## F-statistic: 5.372 on 2 and 16 DF,  p-value: 0.01641
```

Again, it appears not.

Triples

```
triples.sentiments <- read.csv('../data/triples-qualities.csv',
                                header=T,
                                stringsAsFactors = F)

# Get in useful format
triples.sentiments.df <- triples.sentiments %>% group_by(TRIPLE, SENTIMENT) %>%
  summarise(count=n()) %>%
  mutate(total=sum(count),
          prop=count/total)
# remove 'total' variable
triples.sentiments.df$total <- NULL
# Make sure list is complete (some sentiments are missing where none listed in input data)
full_data <- expand.grid(TRIPLE=triples.sentiments.df$TRIPLE, SENTIMENT=triples.sentiments.df$SENTIMENT)
triples.sentiments.df <- unique(left_join(tbl_df(full_data),triples.sentiments.df))

## Joining, by = c("TRIPLE", "SENTIMENT")

## Warning: Column `TRIPLE` joining factor and character vector, coercing into
## character vector

## Warning: Column `SENTIMENT` joining factor and character vector, coercing into
## character vector

triples.sentiments.df$count[is.na(triples.sentiments.df$count)] <- 0
triples.sentiments.df$prop[is.na(triples.sentiments.df$prop)] <- 0

# Order planets and sentiments
triples.sentiments.df$body.1 <- apply(stringr::str_split(triples.sentiments.df$TRIPLE, pattern=" "),
                                     function(x) x[1])
triples.sentiments.df$body.1 <- ordered(ORDER_OF_BODIES[triples.sentiments.df$body.1],
                                     levels=ORDER_OF_BODIES)
triples.sentiments.df$body.2 <- apply(stringr::str_split(triples.sentiments.df$TRIPLE, pattern=" "),
                                     function(x) x[2])
triples.sentiments.df$body.2 <- ordered(ORDER_OF_BODIES[triples.sentiments.df$body.2],
                                     levels=ORDER_OF_BODIES)
triples.sentiments.df$body.3 <- apply(stringr::str_split(triples.sentiments.df$TRIPLE, pattern=" "),
                                     function(x) x[3])
triples.sentiments.df$body.3 <- ordered(ORDER_OF_BODIES[triples.sentiments.df$body.3],
                                     levels=ORDER_OF_BODIES)
triples.sentiments.df$bodies.sorted <- apply(1:nrow(triples.sentiments.df),
                                             function(x)
                                              paste(as.character(sort(unlist(c(triples.sentiments.df[x,"body.1"],
                                                                                   triples.sentiments.df[x,"body.2"],
                                                                                   triples.sentiments.df[x,"body.3"])))
                                                                                   collapse=" "))

# Make body 1 and 2 be in order of bodies as expected
triples.sentiments.df$body.1 <- apply(stringr::str_split(triples.sentiments.df$bodies.sorted, pattern=" "),
                                     function(x) x[1])
triples.sentiments.df$body.2 <- apply(stringr::str_split(triples.sentiments.df$bodies.sorted, pattern=" "),
                                     function(x) x[2])
triples.sentiments.df$body.3 <- apply(stringr::str_split(triples.sentiments.df$bodies.sorted, pattern=" "),
                                     function(x) x[3])
triples.sentiments.df$body.1 <- ordered(triples.sentiments.df$body.1,
                                     levels=ORDER_OF_BODIES)
```

```
triples.sentiments.df$body.2 <- ordered(triples.sentiments.df$body.2,
                                       levels=ORDER_OF_BODIES)
triples.sentiments.df$body.3 <- ordered(triples.sentiments.df$body.3,
                                       levels=ORDER_OF_BODIES)

triples.sentiments.df$order.body.string <- paste0(as.numeric(triples.sentiments.df$body.1),
                                                  as.numeric(triples.sentiments.df$body.2),
                                                  as.numeric(triples.sentiments.df$body.3))
# Order bodies again (this is hacky but it works)
triples.sentiments.df$bodies.sorted <- ordered(triples.sentiments.df$bodies.sorted,
                                              levels=unique(triples.sentiments.df$bodies.sorted[order(t

triples.sentiments.df$SENTIMENT <- ordered(triples.sentiments.df$SENTIMENT,
                                           levels=c("bad", "neutral", "good"))

# Also make a dataframe with an overall sentiment index
triples.overall.sentiments <- triples.sentiments.df %>% group_by(body.1, body.2, body.3, bodies.sorted,
  summarise(sentiment=prop[which(SENTIMENT=="good")]-prop[which(SENTIMENT=="bad")])
```

Now we see how this correlates with component singles or doubles.

```
getDoubleScoresForTriple <- function(triple, double.data=doubles.overall.sentiments, mean=TRUE){
  triple.bodies <- ordered(unlist(stringr::str_split(triple, pattern=" ")),
                          levels=ORDER_OF_BODIES)

  # combinations
  doubles <- ordered(combn(triple.bodies, 2), levels=ORDER_OF_BODIES)
  doubles <- sapply(c(1, 3, 5), function(x) paste(c(doubles[x], doubles[x+1]), collapse= ' '))

  doubles.ordered <- sapply( doubles,
                            function(x) paste(as.character(sort(unlist(strsplit(fixed = TRUE, split=" ", x))), c
  doubles.ordered <- as.data.frame(doubles.ordered)
  set.of.doubles <- sapply(c(1,2, 3),
                          function(x) paste(ORDER_OF_BODIES[as.numeric(as.character(doubles.ordered[1,x]))],
                                              ORDER_OF_BODIES[as.numeric(as.character(doubles.ordered[2,x]))]))

  prop <- 0
  i <- 0
  props <- c()
  for (double in set.of.doubles){
    if (!double %in% c("Mars Sun", "Mars Moon")){
      new.prop <- double.data[which(double.data$bodies.sorted==double), "sentiment"]
      prop <- prop + new.prop
      props <- c(props, as.numeric(new.prop))
      i <- i+1
    }
  }

  mean.prop <- prop/i
  if (mean==TRUE){
    return(as.numeric(mean.prop))
  }
  else{
    return(sum(props))
  }
}
```

```

}
}
triples.overall.sentiments$mean.double.sentiment <- sapply(triples.overall.sentiments$bodies.sorted, fun

```

Now we plot this relationship.


```

spearman.cor.triple.double <- cor.test(triples.overall.sentiments$mean.double.sentiment, triples.overall

## Warning in cor.test.default(triples.overall.sentiments$mean.double.sentiment, :
## Cannot compute exact p-value with ties

p.sentiments.triples.doubles <- ggplot(triples.overall.sentiments, aes(mean.double.sentiment, sentiment
  stat_smooth(method="lm", se=FALSE, colour="red")+
  geom_point(colour="black")+
  theme_basic()+
  ggrepel::geom_text_repel(aes(label=bodies.sorted), size=2)+
  coord_fixed()+
  xlim(c(-1, 1))+
  ylim(c(-1,1))+
  geom_hline(yintercept = 0, linetype='dashed')+
  geom_vline(xintercept = 0, linetype='dashed')+
  xlab("Mean sentiment index of component doubles")+
  ylab("Sentiment index of triple")+
  annotate(geom="label", label=paste0("Spearman's rho = ", myround(spearman.cor.triple.double$estimate,
    "p = ", myround(spearman.cor.triple.double$p.value, 3)), x=1, y=-0
p.sentiments.triples.doubles

```



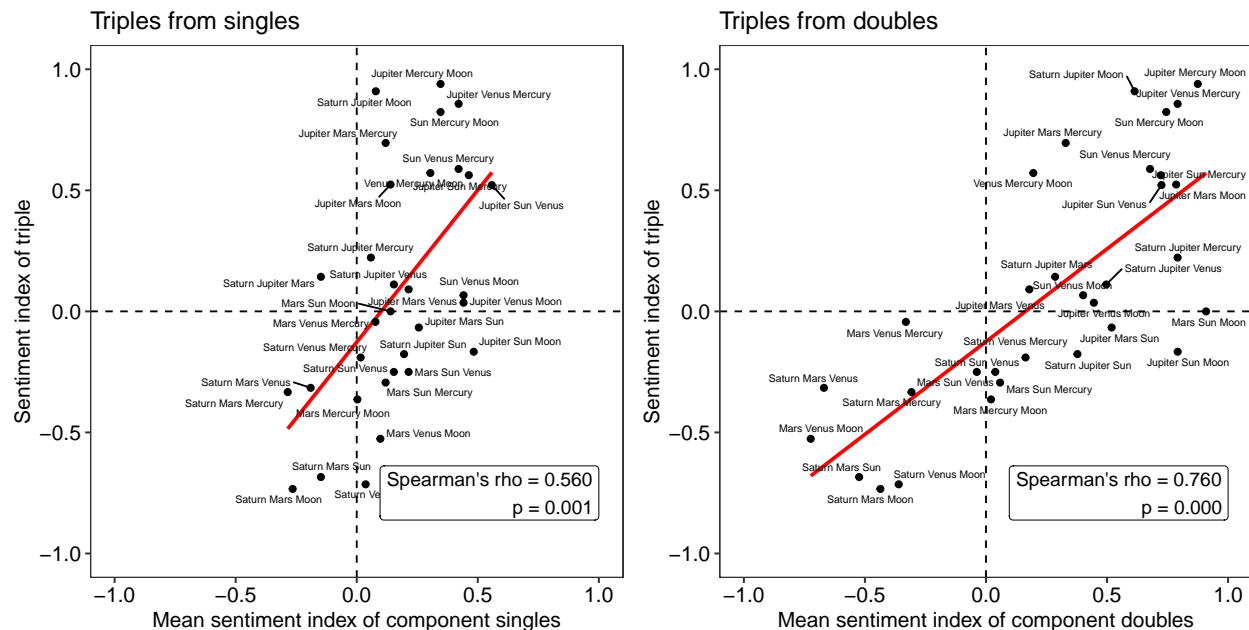
Is this stronger or weaker if we just consider component singles?

```
triples.overall.sentiments$mean.single.sentiment <- sapply(triples.overall.sentiments$bodies.sorted, fun
spearman.cor.triple.single <- cor.test(triples.overall.sentiments$mean.single.sentiment, triples.overall.

## Warning in cor.test.default(triples.overall.sentiments$mean.single.sentiment, :
## Cannot compute exact p-value with ties

p.sentiments.triples.singles <- ggplot(triples.overall.sentiments, aes(mean.single.sentiment, sentiment
  stat_smooth(method="lm", se=FALSE, colour="red")+
  geom_point(colour="black")+
  theme_basic()+
  ggrepel::geom_text_repel(aes(label=bodies.sorted), size=2)+
  coord_fixed()+
  xlim(c(-1, 1))+
  ylim(c(-1,1))+
  geom_hline(yintercept = 0, linetype='dashed')+
  geom_vline(xintercept = 0, linetype='dashed')+
  xlab("Mean sentiment index of component singles")+
  ylab("Sentiment index of triple")+
  annotate(geom="label", label=paste0("Spearman's rho = ", myround(spearman.cor.triple.single$estimate,
                                                                    "p = ", myround(spearman.cor.triple.single$p.value, 3)), x=1, y=-0.5)

# Plot both of them
cowplot::plot_grid(p.sentiments.triples.singles+ggtitle("Triples from singles"), p.sentiments.triples.d
```



Could we model the single sentiments by the occurrence of the double conjunctions in some way? Can we use just rank of median occurrence?

```
doubles.overall.sentiments$rank <- rank(doubles.overall.sentiments$median.occurrence)
summary(lm(sentiment ~ mean.single.sentiments + rank, data=doubles.overall.sentiments))
```

```
##
## Call:
## lm(formula = sentiment ~ mean.single.sentiments + rank, data = doubles.overall.sentiments)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```



```
## -1.1703 -0.2699 0.1485 0.3039 0.7842
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.27050     0.25888   1.045  0.31161
## mean.single.sentiments 1.37606     0.40095   3.432  0.00342 **
## rank             -0.02276     0.02223  -1.024  0.32119
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.527 on 16 degrees of freedom
## Multiple R-squared:  0.4383, Adjusted R-squared:  0.3681
## F-statistic: 6.243 on 2 and 16 DF,  p-value: 0.009908
```

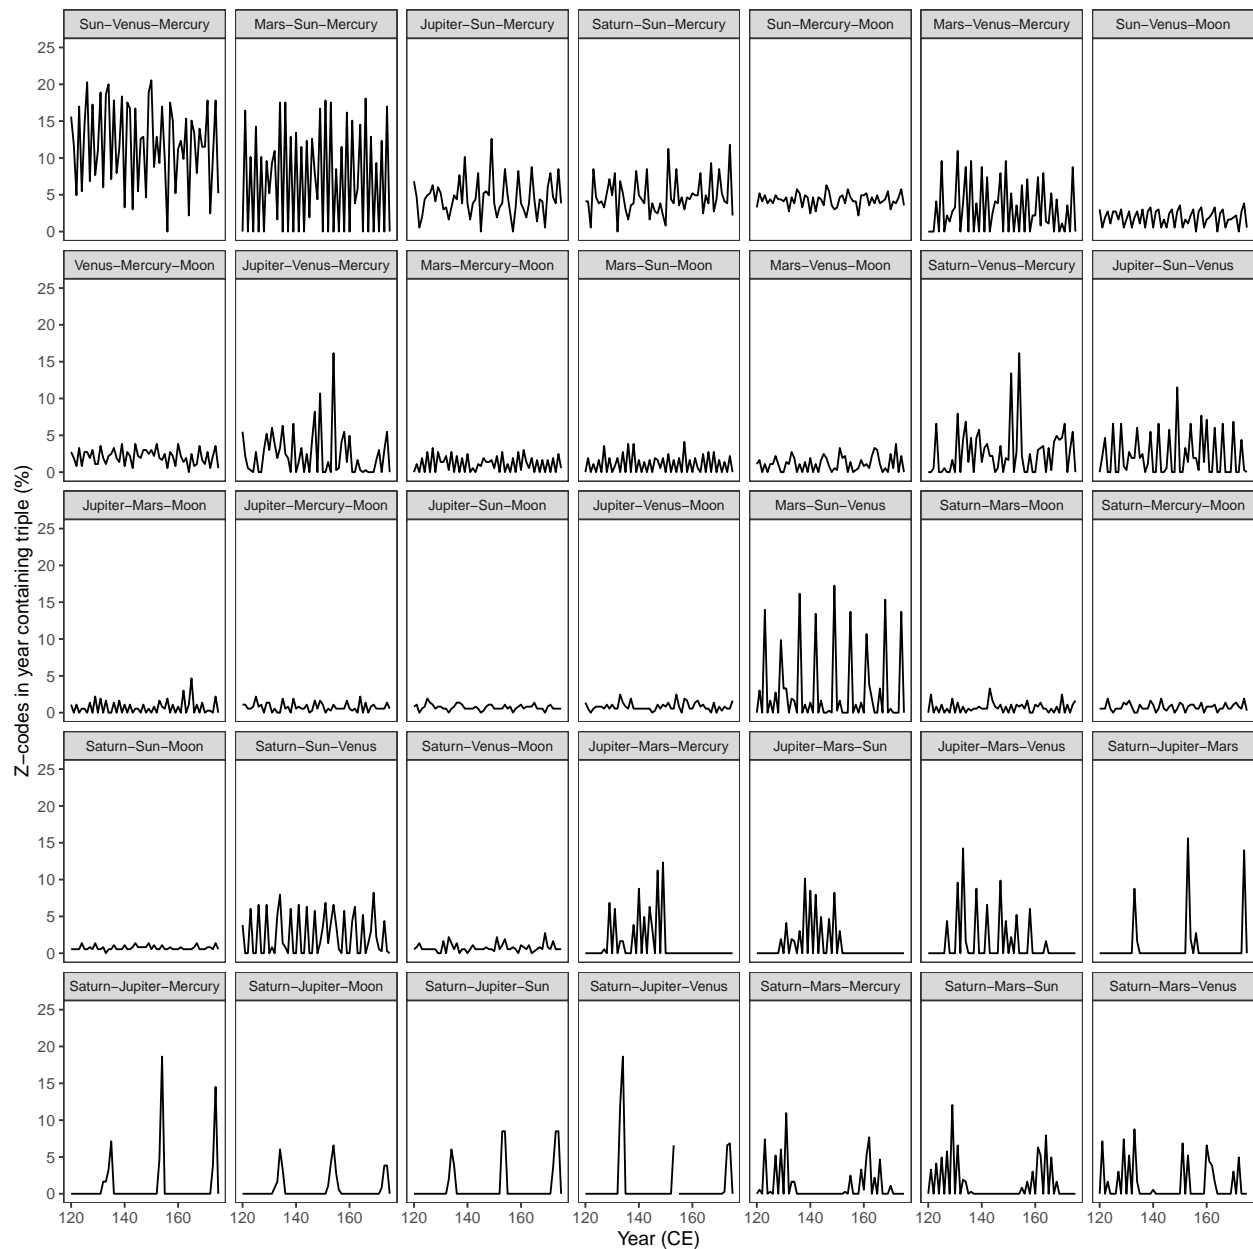
Look at triples occurrences.

```
triple.occurrences <- read.csv('../data/0-CE-200-CE-triple-occurrence-per-year.csv', header=F, stringAsFactors=F)
colnames(triple.occurrences) <- c("year", "triple", "occurrence")
triple.occurrences.median <- triple.occurrences %>% group_by(triple) %>%
  summarise(median=median(occurrence))
order.of.triples <- triple.occurrences.median$triple[order(triple.occurrences.median$median, decreasing=T)]
triple.occurrences.median$year <- 160
triple.occurrences.median$y <- 11
triple.occurrences.median$triple <- ordered(triple.occurrences.median$triple,
  levels=order.of.triples)
triple.occurrences.median$median.plot <- myround(triple.occurrences.median$median, 1)
triple.occurrences.median$triple <- ordered(triple.occurrences.median$triple,
  levels=order.of.triples)
triple.occurrences$triple <- ordered(triple.occurrences$triple,
  levels=order.of.triples)

# Make plot
occurrences.plot.triples <- ggplot(triple.occurrences, aes(year, occurrence, group=triple))+
  geom_line()+
  theme_bw()+
  xlab("Year (CE)") +
  ylab("Z-codes in year containing triple (%)") +
  facet_wrap(~triple, ncol=7) +
  theme(panel.grid = element_blank()) +
  ylim(c(0,25)) +
  xlim(c(120, 175)) +
  theme(strip.text = element_text(size=8)) +
  geom_text(data=triple.occurrences.median,
    aes(year, 90, label=paste0(median.plot, "%"), group=triple),
    hjust=0, size=3)
occurrences.plot.triples
```

```
## Warning: Removed 5005 rows containing missing values (geom_path).
```

```
## Warning: Removed 35 rows containing missing values (geom_text).
```



Textual analysis

Using tidytext to look computationally at similarity of descriptions.

First, consider just the length of the descriptions.

Length of descriptions play a role?

```
getLengthPassage <- function(conjunction){
  df <- tibble(word=scan(paste0('../ ../ ../ Texts for analysis/names-removed-Riley English Translation (2017)
  mutate(line=row_number())
  return(as.numeric(nrow(df)))
}
```


Add to doubles

```
doubles.overall.sentiments$length.passage <- as.numeric(sapply(gsub(" ", "-", doubles.overall.sentiment
```

```

    function(x) getLengthPassage(x)))
# Plot
ggplot(doubles.overall.sentiments, aes(median.occurrence, length.passage)) +
  geom_point()+
  ggrepel::geom_text_repel(aes(label=bodies.sorted), size=2)

```



```

# Note that the Sun and the Moon have the longest description

# Are they explained by length of the singles?
singles.overall.sentiments$length.passage <- as.numeric(sapply(gsub(" ", "- ", singles.overall.sentiment.
    function(x) getLengthPassage(x)))
getSingleLengths <- function(conjunction, mean=TRUE, singles=singles.overall.sentiments){
  conjunction.bodies <- ordered(unlist(stringr::str_split(conjunction, pattern=" ")),
    levels=ORDER_OF_BODIES)

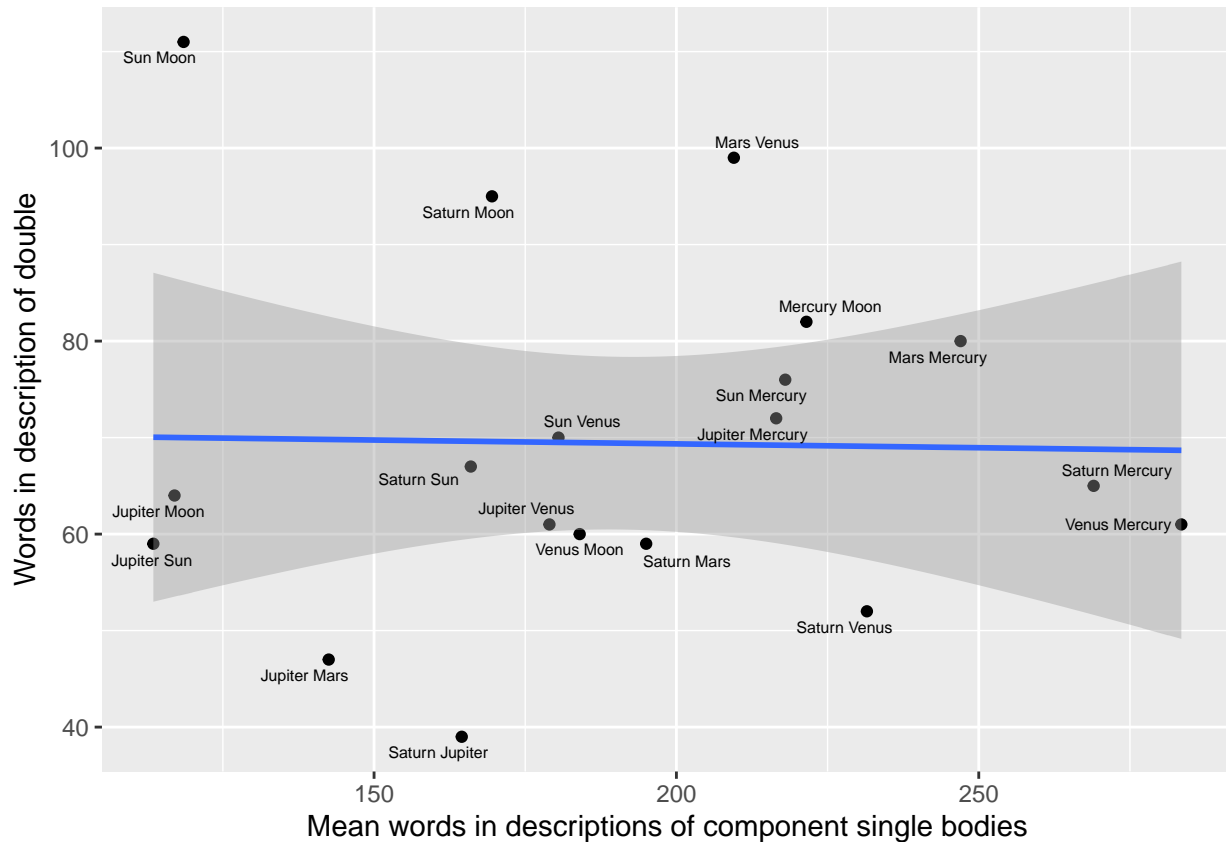
  prop <- 0
  for (body in conjunction.bodies){
    new.prop <- as.numeric(singles[which(singles$PLANET==body), "length.passage"])
    prop <- prop + new.prop
  }
  if (mean==FALSE){
    return(prop)
  }
  else{
    return(prop/length(conjunction.bodies))
  }
}
doubles.overall.sentiments$mean.length.single <- sapply(doubles.overall.sentiments$bodies.sorted,

```

```

function(x) getSingleLengths(x))
ggplot(doubles.overall.sentiments, aes(mean.length.single, length.passage))+
  geom_point()+stat_smooth(method="lm")+
  xlab("Mean words in descriptions of component single bodies")+
  ylab("Words in description of double")+
  ggrepel::geom_text_repel(aes(label=bodies.sorted), size=2)

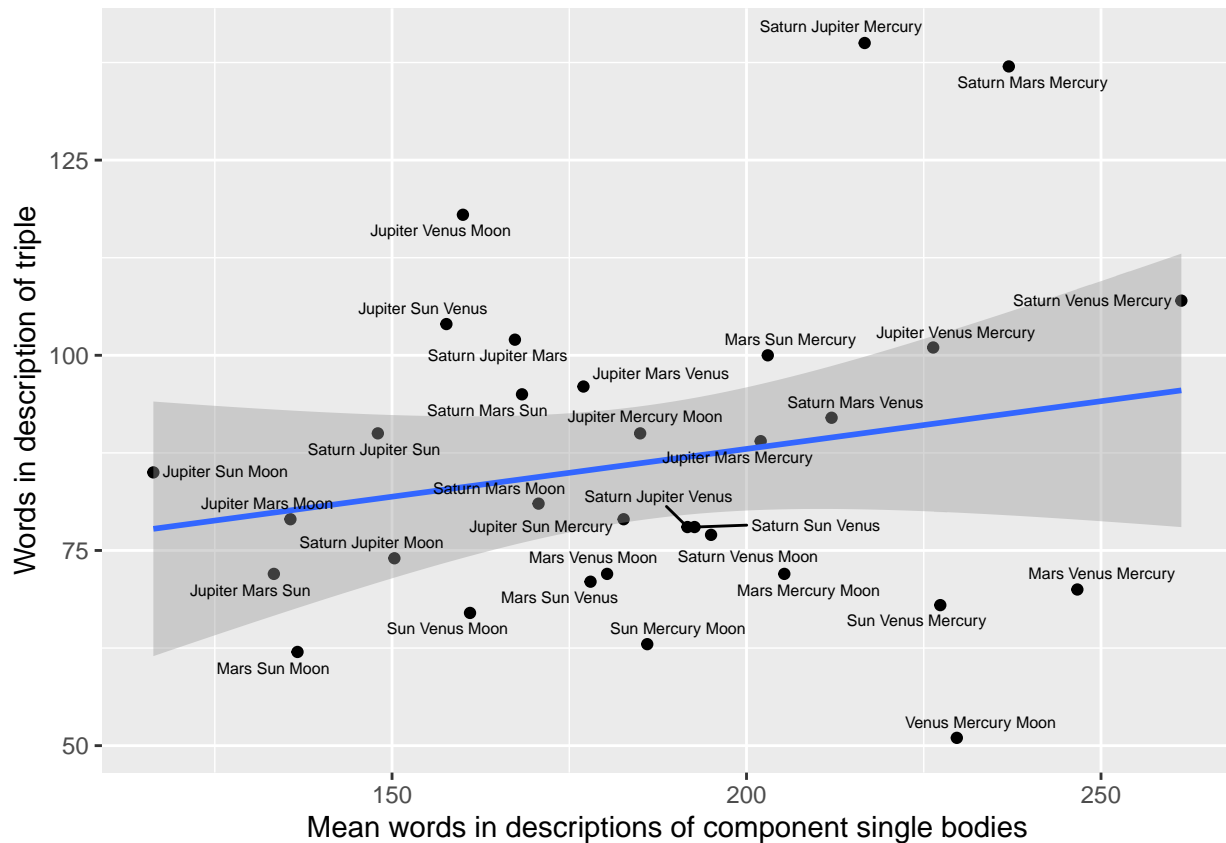
```



```

# Triples lengths
triples.overall.sentiments$length.passage <- as.numeric(sapply(gsub(" ", "-", triples.overall.sentiment.
  function(x) getLengthPassage(x)))
triples.overall.sentiments$mean.length.single <- sapply(triples.overall.sentiments$bodies.sorted,
  function(x) getSingleLengths(x))
# Are these explained by the lengths of the singles?
ggplot(triples.overall.sentiments, aes(mean.length.single, length.passage))+
  geom_point()+stat_smooth(method="lm")+
  xlab("Mean words in descriptions of component single bodies")+
  ylab("Words in description of triple")+
  ggrepel::geom_text_repel(aes(label=bodies.sorted), size=2)

```



Weak positive relationship, heavily driven by the outliers: Saturn Jupiter Mercury and Saturn Mars Me

```
library(tidytext)
conjunction <- "Jupiter Mars Mercury"
df1 <- tibble(word=scan(paste0('.../.../Texts for analysis/names-removed-Riley English Translation (20
  mutate(line=row_number(), name=conjunction)
conjunction <- "Jupiter Mars Venus"
df2 <- tibble(word=scan(paste0('.../.../Texts for analysis/names-removed-Riley English Translation (20
  mutate(line=row_number(), name=conjunction)
for (conjunction in as.character(triples.overall.sentiments$bodies.sorted)){
  df <- tibble(word=scan(paste0('.../.../Texts for analysis/names-removed-Riley English Translation (20
  mutate(line=row_number(), conjunction=conjunction)
  if (exists("full.df")){
    full.df <- rbind(full.df, df)
  }
  else{
    full.df <- df
  }
}

# Co-occurring words
library(widyr)
library(stopwords)
full.df.stop <- full.df %>%
  anti_join(get_stopwords())
```

```
## Joining, by = "word"
```

```

full.df.stop <- full.df.stop %>%
  filter(!word %in% c("They", "men", "make", "bring", "cause", "These", "stars", "become", "men,"))
words_by_conjunction <- full.df.stop %>%
  pairwise_count(item=word, feature=conjunction, sort=TRUE)
# Words that appear in the same overall description
full.df.stop.sorted <- full.df.stop %>%
  count(word, sort=TRUE)

# Similarity
for (conjunction in as.character(doubles.overall.sentiments$bodies.sorted)){
  df <- tibble(word=scan(paste0('../../Texts for analysis/names-removed-Riley English Translation (2
  mutate(line=row_number(), conjunction=conjunction)
  if (exists("doubles.descriptions")){
    doubles.descriptions <- rbind(doubles.descriptions, df)
  }
  else{
    doubles.descriptions <- df
  }
}
for (planet in as.character(ORDER_OF_BODIES)){
  df <- tibble(word=scan(paste0('../../Texts for analysis/names-removed-Riley English Translation (2
  mutate(line=row_number(), conjunction=planet)
  if (exists("single.descriptions")){
    single.descriptions <- rbind(single.descriptions, df)
  }
  else{
    single.descriptions <- df
  }
}

# Look at Saturn to start with
doubles.descriptions.saturn <- doubles.descriptions %>% mutate(Saturn=grep1("Saturn", conjunction))
# are those with saturn more similar to saturn to those without saturn
saturn.words <- (doubles.descriptions.saturn %>% filter(Saturn==TRUE))$word
nonsaturn.words <- (doubles.descriptions.saturn %>% filter(Saturn==FALSE))$word
table(saturn.words %in% (single.descriptions %>% filter(conjunction=="Saturn"))$word)

##
## FALSE TRUE
## 238 139

table(saturn.words %in% (single.descriptions %>% filter(conjunction!="Saturn"))$word)

##
## FALSE TRUE
## 217 160

table(sample(nonsaturn.words, size=length(saturn.words)) %in% (single.descriptions %>% filter(conjuncti

##
## FALSE TRUE
## 232 145

table(sample(nonsaturn.words, size=length(saturn.words)) %in% (single.descriptions %>% filter(conjuncti

##
## FALSE TRUE

```

```
##      192      185
table(sample(saturn.words, size=100) %in% sample((single.descriptions %>% filter(conjunction=="Mercury"))

##
## FALSE  TRUE
##      76      24

# Get positive sentiments from Bing 2004 corpus
positive <- get_sentiments("bing") %>% filter(sentiment == "positive")
# Look at positive words
df %>%
semi_join(positive) %>%
count(word, sort = TRUE)

## Joining, by = "word"

## # A tibble: 0 x 2
## # ... with 2 variables: word <chr>, n <int>
```